

Supervised Machine Learning I: Prediction

Paul Goldsmith-Pinkham

April 29, 2021

Machine Learning + Causality

- Today, focusing on how to tie machine learning methods into estimation of causal effects
- Most of our ideas revolve around how to think about estimating CATEs – conditional average treatment effects
 - Why is this interesting? Why is knowing CATEs preferable to ATEs?
- Recall that with exhaustively defined CATEs, we can estimate our ATE
 - But, crucially, we could *target* appropriately
 - Well-estimated CATEs help identify better decisions based on decision rules
 - Also good for economic theory!
- But, can be hard to do in a disciplined way

Why can ML be powerful in this space?

- A serious concern in empirical work is specification hunting – looking for significant effects in subgroups, and then telling a story about it
- One solution is pre-analysis plans – tying our hands before the fact about what we will look at
- However, sometimes we would like to let the “data speak”
 - What if we could automate the process for estimating significant CATEs?
- Machine learning could allow us to estimate these approaches in a standardized way, while using out-of-sample testing to ensure that we are not data mining

The literatures with Machine learning and CATEs



- Today, will talk about two papers/lits:
 - Causal Trees (From Athey and Imbens (2016)
 - More generally in the space of causal partitioning
 - Causal “Forests” by Wager and Athey (2019) as well
 - GATES and CLAN from Chernozhukov et al. (2020)
 - GATES = Sorted Group Average Treatment Effects
 - CLAN = Classification Analysis
- These approaches are similarly focusing on CATEs, but solving a crucial statistical problem in two distinct ways

Machine learning and CATEs

- What is the statistical problem? There are two (related) issues:
 1. Inference: even if we predict or show the effect of a treatment is higher in one subgroup than another, can we say whether this is just due to random variation, or a meaningful difference?
 2. Testing causal inference out-of-sample: Evaluating how “accurate” you are requires knowing your target outcome. E.g. $Y_i - \hat{Y}_i$. But, $\tau_i = Y_i(1) - Y_i(0)$ is fundamentally unknown.
- These issues are in large part solved by additional sample splitting
- Importantly: these approaches do *not* solve the issue of exogeneous variation
 - In most settings, this should be viewed as a setting where we have a randomly varying treatment (e.g. an RCT) and we want to study CATEs
 - However, if we have a good IV, we could study the reduced form quite sensibly!



Causal trees (Athey and Imbens (2016))

- Necessary notation: Binary treatment, D_i , and covariates (potentially high dimensional) X_i . Outcome Y_i .
- In our discussion, we'll assume completely random assignment of D_i , but it is possible to account for conditioning variables as well using a p-score method
- The key approach will be following the tree-based approach from last class, but with some essential modifications
 - Recall that trees worked by splitting up observations at a given node based on a given characteristic



Causal trees (Athey and Imbens (2016))

- Key insight of this paper: when you choose what to split on, you are picking something that is “unusual” relative to the underlying data generating process
 - This induces bias!
- Hence, the CT approach splits the sample again – first using part of the sample to pick the tree leaves, then testing the effects within the leaves using the left-out sample
 - This gives you three samples: two training, and one true test



Causal trees (Athey and Imbens (2016))

- How to implement? In R, there's a very nice package that includes the Random Forest (see Wager and Athey (2019)) approach as well:

<https://grf-labs.github.io/grf/>

The GRF Algorithm

The following guide gives an introduction to the generalized random forests algorithm as implemented in the `grf` package. It aims to give a complete description of the training and prediction procedures, as well as the options available for tuning. This guide is intended as an informal and practical reference; for a theoretical treatment of GRF, please consult the 'Generalized Random Forests' paper.



GRF extends the idea of a classic random forest to allow for estimating other statistical quantities besides the expected outcome. Each forest type, for example `quantile_forest`, trains a random forest targeted at a particular problem, like quantile estimation. The most common use of GRF is in estimating treatment effects through the function `causal_forest`.

Causal trees (Athey and Imbens (2016))

- How to implement? In R, there's a very nice package that includes the Random Forest (see Wager and Athey (2019)) approach as well:

<https://grf-labs.github.io/grf/>

- For Python, the `econml` package can do this as well (as well as many other approaches): <https://econml.azurewebsites.net/spec/spec.html>

Welcome to `econml`'s documentation! 🔗

- [EconML User Guide](#)
 - [Machine Learning Based Estimation of Heterogeneous Treatment Effects](#)
 - [Motivating Examples](#)
 - [Customer Targeting](#)
 - [Personalized Pricing](#)
 - [Stratification in Clinical Trials](#)
 - [Learning Click-Through-Rates](#)
 - [Problem Setup and API Design](#)
 - [API of Conditional Average Treatment Effect Package](#)
 - [Linear in Treatment CATE Estimators](#)
 - [Example Use of API](#)

Causal trees (Athey and Imbens (2016))

- How to implement? In R, there's a very nice package that includes the Random Forest (see Wager and Athey (2019)) approach as well:
<https://grf-labs.github.io/grf/>
- For Python, the econml package can do this as well (as well as many other approaches): <https://econml.azurewebsites.net/spec/spec.html>
- Nothing in Stata, sorry



GATES and CLAN


- The causal tree approach is a beautiful approach in solving the bias and inference issues
- However, the general inference solution does not account for the uncertainty in the binning of the covariates
 - Recall how the method works – by using a split sample to choose the bins, the CATEs within those bins work just as well as any standard regression approach
 - But this fails to account for the fact that these bins may change in different samples
- Chernozhukov et al. (2020) highlight this issue, and propose a much more general approach
 - This approach has more limitations, but at the benefit of being even more general

GATES and CLAN

- This approach has a lot of technical details
 - I will not be able to do it justice today
- However, the key concept, as highlighted in the paper, is that instead of trying to identify the CATEs directly, identify the key features of the CATEs instead
 - More precisely, identify how much heterogeneity there is in the underlying estimates
 - Then, figure out the characteristics of those groups with heterogeneous effects
- The key approach starts with the following concept:
 - Randomly split the sample into a main and auxiliary sample
 - In the auxiliary sample, estimate the control mean, $B(X)$ and the treatment effect $\tau(X)$ using any ML method.
 - With these predictions, we will now proceed

GATES and CLAN

- The problem is that $\tau(X)$ is very high-dimensional
- The GATES approach says – what if we grouped the effects into bins G , increasing in effect size.
 - We can talk about the property of these GROUPED average treatment effects, which average of the high dimensional properties
 - In turns out we can say a lot about that, statistically
- Moreover, we can test for whether these are all the same
 - Harkens back to binscatter and testing for monotonicity!



GATES and CLAN

- The issue is that we still haven't solved for what these groups are
 - Knowledge of heterogeneity doesn't get us very far
- The CLAN approach asks how important characteristics vary by these binned groups
- We can use this to identify bins worth targetting

TABLE 5. CLAN of Immunization Incentives

	Elastic Net		
	20% Most (δ_5)	20% Least (δ_1)	Difference ($\delta_5 - \delta_1$)
Number of vaccines to pregnant mother	2.161 (2.110,2.212)	2.288 (2.237,2.337)	-0.128 (-0.200,-0.055) [0.001]
Number of vaccines to child since birth	4.230 (4.100,4.369)	4.714 (4.573,4.860)	-0.513 (-0.710,-0.311) [0.000]
Fraction of children received polio drops	1.000 (1.000,1.000)	1.000 (1.000,1.000)	0.000 (0.000,0.000) [0.000]
Number of polio drops to child	2.964 (2.954,2.975)	2.998 (2.987,3.007)	-0.033 (-0.047,-0.019) [0.000]
Fraction of children received immunization card	0.899 (0.878,0.922)	0.932 (0.908,0.956)	-0.036 (-0.065,-0.004) [0.000]
Fraction of children received Measles vaccine by 15 months of age	0.127 (0.100,0.155)	0.255 (0.230,0.282)	-0.131 (-0.167,-0.094) [0.052]
Fraction of children received Measles at credible locations	0.290 (0.252,0.327)	0.435 (0.400,0.470)	-0.152 (-0.198,-0.097) [0.000]

Notes: Medians over 100 splits. 90% confidence interval in parenthesis.

Notes: P-values for the hypothesis that the parameter is equal to zero in brackets.

Implementation in practice

- Chernozhukov et al. (2020) outline the algorithm in detail in the paper
- Best I have found is Max Eber has code here:
`https://github.com/maximilianeber/ml-treat`
- Otherwise... good luck! I think this is very doable, if you've done it once, but there's a serious learning curve